

CS6023: GPU Programming

Assignment 4 (20 marks)

Deadline : April 24, 2022, 23:55 on Moodle

1 Problem Statement

The aim of this assignment is to write a *Train Ticket Booking* application. There are N trains with the train numbers being $0, 1, 2, \dots, N - 1$. Each train has a variable M number of classes (the classes being $0, 1, 2, \dots, M - 1$) with each class having a maximum capacity. Each train is also associated with a source station (src) and a destination station ($dest$). The src station is **never** the same as the $dest$ station.

The list of station that the train goes through is either $[src, src + 1, src + 2, \dots, dest]$ or $[dest, dest - 1, dest - 2, \dots, src]$ depending on whether $src < dest$ or $src > dest$ respectively.

The booking requests are split across batches and they are processed in the following manner:

- All the requests in Batch i are processed before the requests in Batch $i + 1$ are processed.
- All the requests within the same batch belonging to the same train and same class are to be processed in order of their *RequestID*. (See Input Format below)
- Requests belonging to different trains / same train and different classes are to be processed in parallel so as to maximise the number of requests being processed at a given time.

2 Input and Output

2.1 Input Format

- First line contains N // the number of trains.
- For each of the N trains, the following is given:
 - Train Number, Number of Classes (M), Source (src), Destination ($dest$)
 - The next M lines contains Class Number, Max Capacity
- Next line contains B //the number of batches
- For each of the B batches, the following is given:
 - R //Number of requests in the given batch
 - Each of the R requests is of the form: Request ID, Train Number, Class, Source, Destination, Number of Seats

2.2 Output Format

For each request in the batch, print either *success* or *failure* depending on whether the request could be processed or not. After each batch is processed, the following is to be printed in different lines:

- Number of requests that succeeded <space> Number of requests that failed
- Total number of seats that were booked in the given batch.

The total number of seats that are booked when a particular request is processed can be calculated in the following manner:

- Say, the request is for booking n seats from src to $dest$.
- Then the number of seats booked is $n \times abs(dest - src)$ where abs is the *absolute* function.

2.3 Constraints

- $1 \leq N \leq 100000$ (number of trains)
- $1 \leq M \leq 25$ (number of classes in each train)

- $1 \leq B \leq 5000$ (number of batches of requests)
- $1 \leq R \leq 5000$ (number of requests in a batch)
- $1 \leq C \leq 100$ (max capacity of any class)
- $1 \leq \text{abs}(\text{dest} - \text{src}) \leq 50$
- $\text{src} \geq 0, \text{dest} \geq 0$

3 Sample Testcase

3.1 Sample Input

```

2 //Number of Trains
0 2 57 54 //(Train Num - 0, Num Classes - 2, Src - 57, Dest - 54)
0 10 //(Class - 0, Max Capacity - 10)
1 12 // ...
1 3 45 49 //(Train Num - 1, Num Classes - 3, Src - 45, Dest - 49)
0 12 //(Class - 0, Max Capacity - 12)
1 5 // ...
2 10 // ...
2 //Number of batches of requests
4 //4 requests in 1st batch
0 0 0 56 54 4 //(RequestID, TrainNum,ClassNum,Src,Dest,NumSeats)
1 0 0 56 54 4 // ...
2 0 0 56 54 3 // ...
3 1 2 45 49 12 // ...
1 //1 request in 2nd batch
0 0 0 57 56 5 //(RequestID, TrainNum,ClassNum,Src,Dest,NumSeats)

```

3.2 Sample Output

```

success
success
failure
failure
2 2
16
success
1 0

```

The first three requests in the 1st batch of requests try to book seats in the 0th class of the 0th train. Since the maximum capacity for the mentioned class is only 10, the first and second request succeeds and the third request fails. The fourth request in the 1st batch also fails because the number of seats asked for is 12 whereas the maximum capacity for the asked class is only 10 (train 1, class 2). Thus, in the 1st batch there are 2 requests which succeeds and 2 requests which fails. Hence, 2 2 in the 5th line of the output. The total number of seats booked is 16 $((4 \times 2) + (4 \times 2))$.

There is only 1 request in the 2nd batch and it succeeds because the requirement is satisfied.

4 Points to be noted

- Say, a booking request is from *src* station to *dest* (and $src < dest$) station for n seats succeeds. Then, all the n seats are occupied from station *src* to station $(dest - 1)$ and the seats become free at the *dest* station. For example, if the max capacity for a particular class is 50, *src* station is 0, *dest* station is 10, then both the following requests succeed:
 - 50 seats from 0 to 5.
 - 50 seats from 5 to 10.
- For this assignment, the reading/writing of input/output, use of data structures, kernel prototypes has been completely left to your choice.
- Sequential implementation (apart from where absolutely needed) of the above functionality will lead to '0' marks on the assignment.
- Test your code on large inputs.

5 Submission Guidelines

- Name the file as 'RollNumber.cu' which contains the implementation of the above-described functionality.

- Submit **only** the 'RollNumber.cu' file to Moodle. For example, if your roll number is CS17B049, then the name of the file to be submitted on Moodle is 'CS17B049.cu'.
- Read input from standard input and print the output to standard output so that we can use I/O redirection for the evaluation.
- After submission, download the file and make sure it was the one you intended to submit.
- Non-adherence to the above guidelines **will** lead to a penalty.

6 Learning Suggestions

- Write a CPU-version of code achieving the same functionality. Time the CPU code and GPU code separately for large matrices and compare the performances.
- Exploit shared memory as much as possible to gain performance benefits.